



Concurrent Web Server

Contributors:

Kanishk Kalra (17110067)

Kavita Vaishnaw (17110073)

Mohit Mina (17110078)

Prakash R (17110109)

Shaurya Agarwal (17110145)

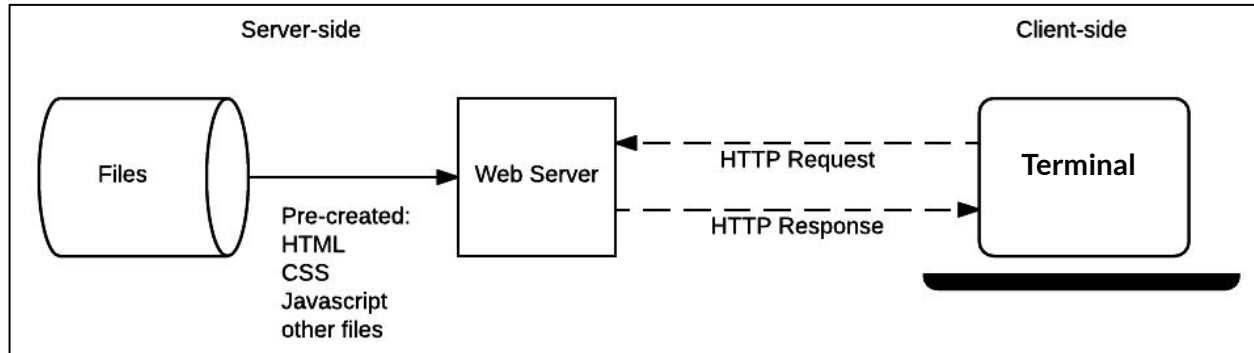


Project Description

Make a concurrent web server with scheduling policies and security considerations.

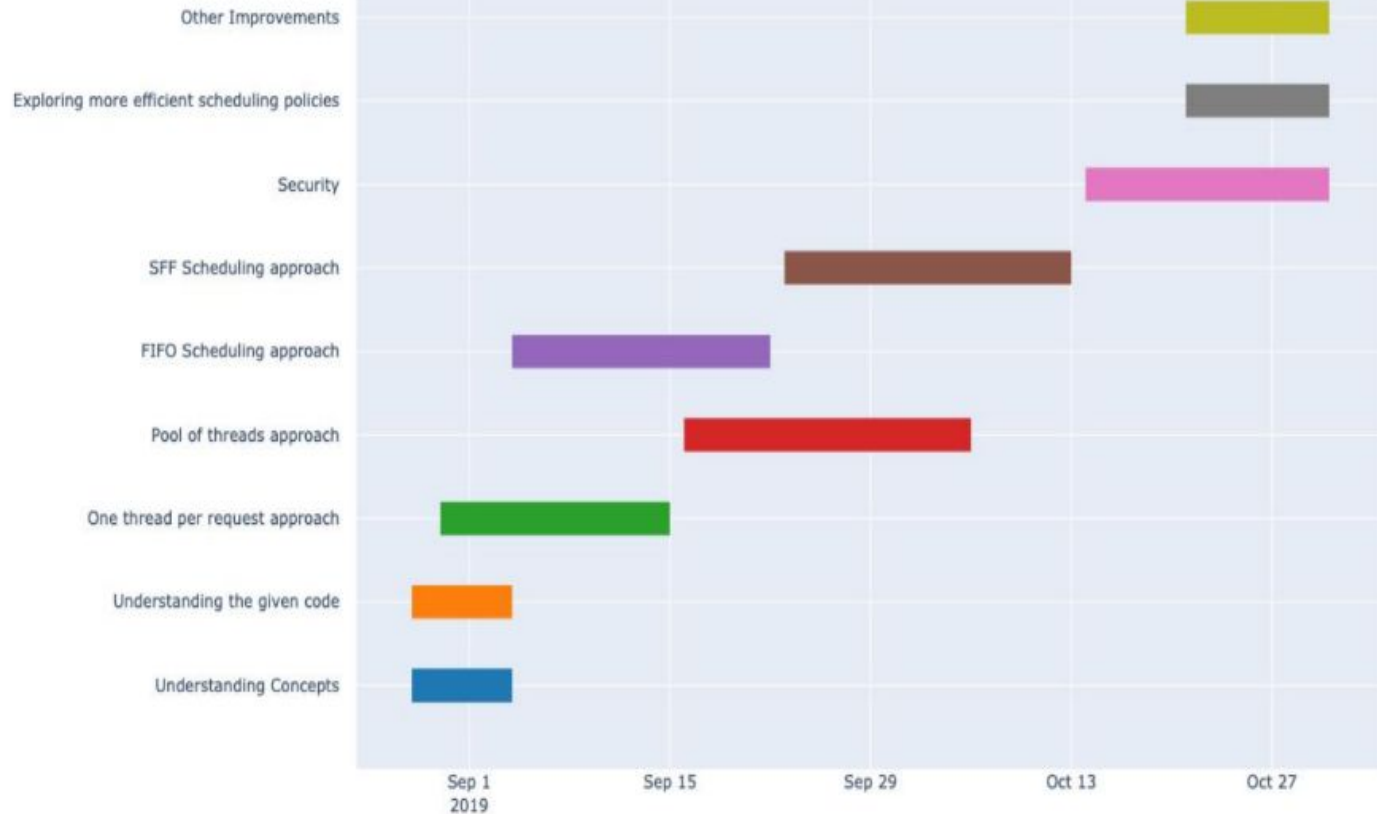
Aim:

- Webserver should handle multiple requests simultaneously.
- It should have 2 scheduling policies (FCFS/FIFO and SFF).
- Web server should have some security considerations..

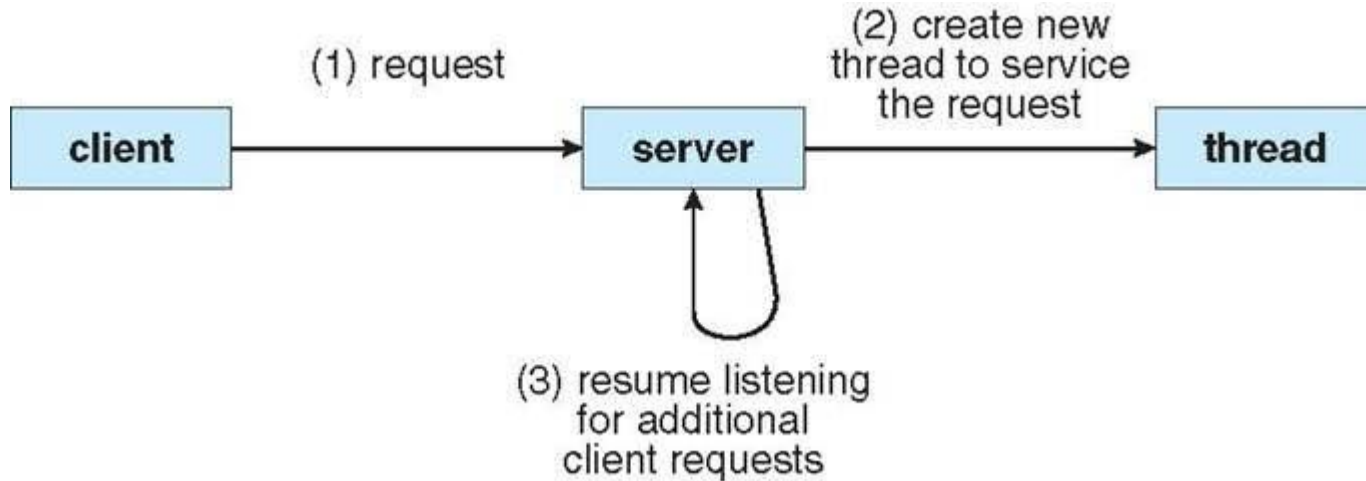


Source: <https://mdn.mozillademos.org/files/13841/Basic%20Static%20App%20Server.png>

Gantt Chart



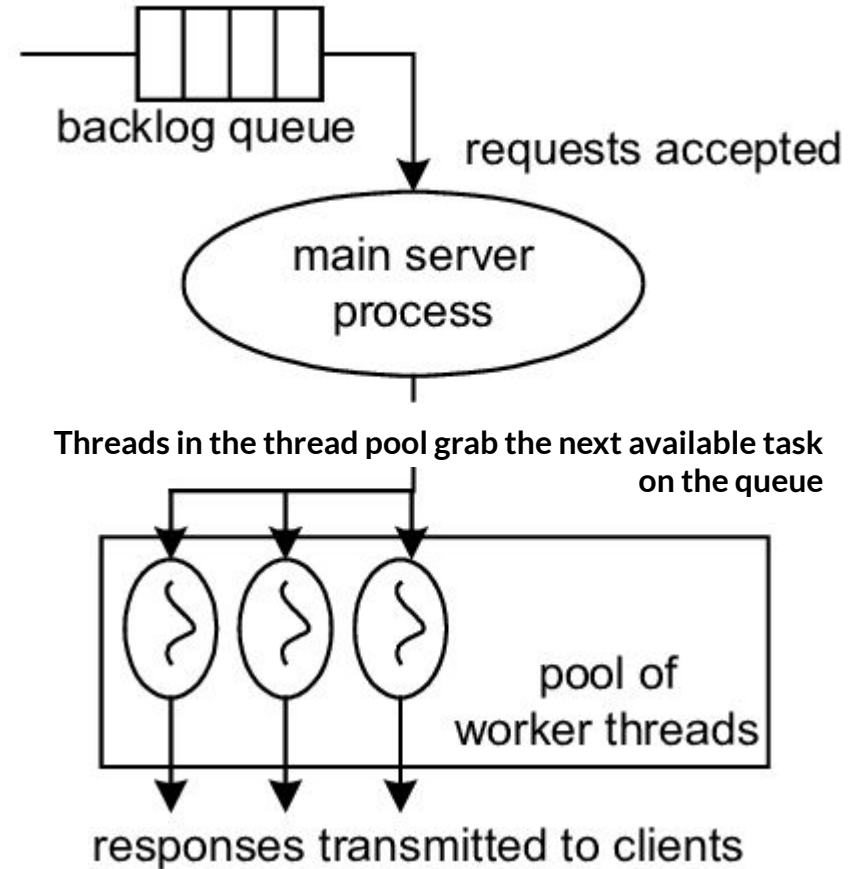
One Thread per Request Approach



Source:

<http://www.shubhsblog.com/programming/multithreaded-webserver-java.html>

Pool of Threads with FCFS Scheduling



Source:

<https://www.classes.cs.uchicago.edu/archive/2016/spring/12300-1/threadpool.png>



Structure of Thread Pool

```
18 struct thread_pool{
19     pthread_t *workers; //worker threads
20     struct q_node* queue; //queue that holds tasks waiting to be executed
21     int q_head, q_tail;
22     int max_workers;
23     int TASK_Q_MAX;
24     int scheduled; //keeps track of number of tasks scheduled
25     pthread_mutex_t mutex;
26     pthread_cond_t task_available; //signaled on when we go from no task to task available
27     pthread_cond_t no_task_left; //signaled on when no more tasks scheduled
28
29 };
```

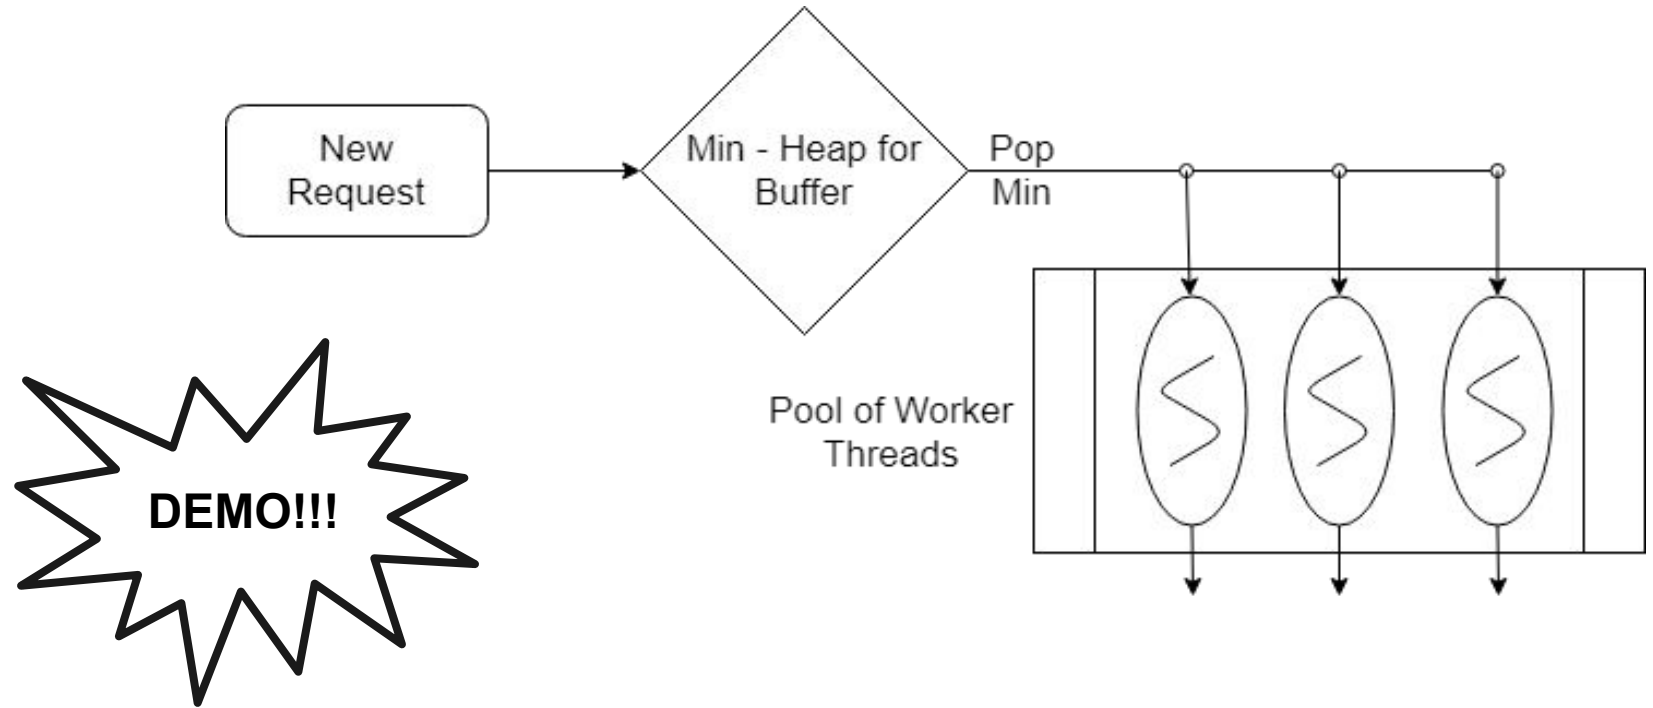
```

49 void *worker_func(void *pool_arg){
50     struct thread_pool *pool = (struct thread_pool *)pool_arg;
51     while(1){
52         struct q_node task_picked;
53         pthread_mutex_lock(&pool->mutex);
54
55         while (pool->q_head == pool->q_tail){
56             //empty queue
57             pthread_cond_wait(&pool->task_available, &pool->mutex);
58         }
59
60         //assert(pool->q_head != pool->q_tail);
61         task_picked = pool->queue[pool->q_head % pool->TASK_Q_MAX];
62         pool->q_head++;
63         pool->scheduled++; //task scheduled
64
65         pthread_mutex_unlock(&pool->mutex);
66
67         task_picked.routine(task_picked.arg);
68
69         pthread_mutex_lock(&pool->mutex);
70
71
72
73         pool->scheduled--;
74
75         if (pool->scheduled == 0){
76             pthread_cond_signal(&pool->no_task_left);
77         }
78
79         pthread_mutex_unlock(&pool->mutex);
80     }
81
82     return NULL;
83 }

```

Worker Function

Shortest File First (SFF) Scheduling




```

235 ▾ void pool_add_task(struct thread_pool *pool, void *(*routine)(void*), void *arg, void *tem){
236     struct info* temp = (struct info*) tem;
237     pthread_mutex_lock(&pool->mutex);
238     if ((pool->heap)->count==0){
239         pthread_cond_broadcast(&pool->task_available);
240     }
241     struct q_node task;
242     task.routine = routine;
243     // task.arg = arg;
244     task.size = temp->size;
245     task.req = *temp;
246     insert(pool->heap, task);
247     print_heap(pool);
248     pthread_mutex_unlock(&pool->mutex);
249 }
250





```

```

124 ▾ void insert(Heap *h, struct q_node key){
125 ▾     if( h->count < h->capacity){
126         h->arr[h->count] = key;
127         heapify_bottom_top(h, h->count);
128         h->count++;
129     }
130 }
131

```

Security



Apache Benchmark Testing

Apache Bench (ab) is a load testing and benchmarking tool for Hypertext Transfer Protocol (HTTP) server.

```
Server Software:      OSTEP
Server Hostname:      localhost
Server Port:          10000

Document Path:        /
Document Length:      3271 bytes

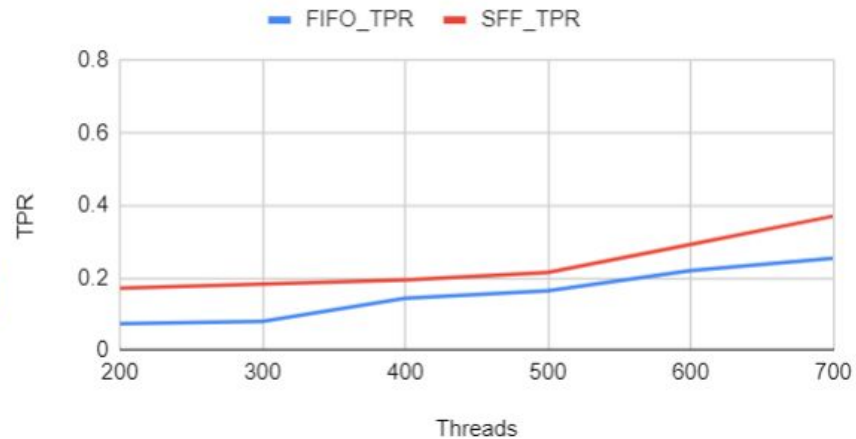
Concurrency Level:     100
Time taken for tests:  0.007 seconds
Complete requests:     100
Failed requests:       0
Total transferred:     336200 bytes
HTML transferred:     327100 bytes
Requests per second:   13548.30 [#/sec] (mean)
Time per request:      7.381 [ms] (mean)
Time per request:      0.074 [ms] (mean, across all concurrent requests)
Transfer rate:         44481.82 [Kbytes/sec] received

Connection Times (ms)
              min    mean[+/-sd] median    max
Connect:        1      2    0.3      2      2
Processing:      1      2    1.3      2      4
Waiting:         1      2    1.2      1      4
Total:          3      4    1.0      3      5


Percentage of the requests served within a certain time (ms)
 50%      3
 66%      4
 75%      5
 80%      5
 90%      5
 95%      5
 98%      5
 99%      5
100%      5 (longest request)
```

TPR Comparison

TPR (Time Per Request) -n 300 -c 50

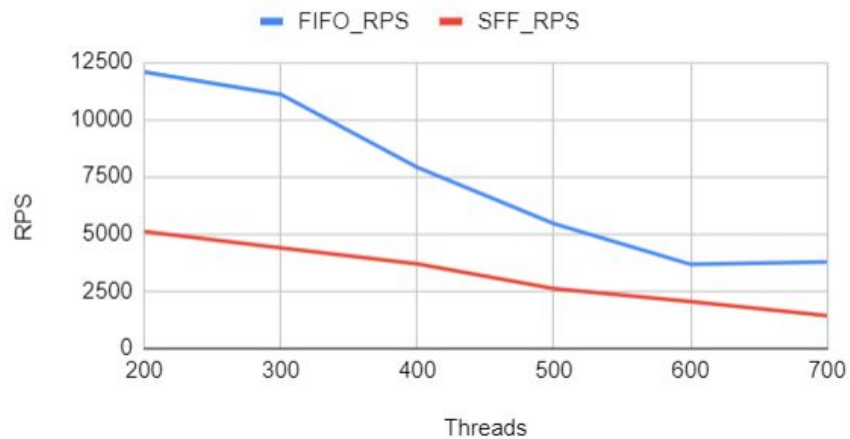


TPR (Time Per Request) -n 300 -c 25

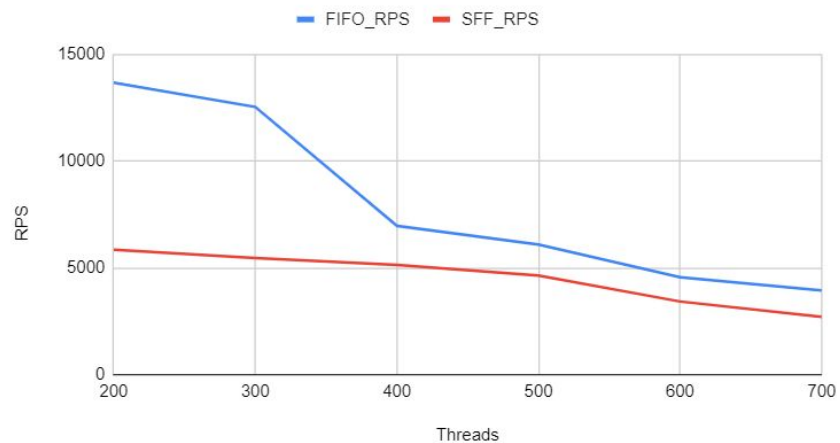


RPS Comparison

RPS (Requests Per Second) -n 300 -c 25



RPS (Requests Per Second) -n 300 -c 50



Contributions of Members



- ★ One Thread Per Request Approach: Kanishk, Shaurya
- ★ Web Client: Mohit, Kanishk, Shaurya
- ★ Pool of Threads with FCFS Scheduling: Prakash, Kavita
- ★ SFF Scheduling: Prakash, Kavita
- ★ Security: Kanishk, Shaurya
- ★ Apache Benchmarking: Mohit

References



- <https://github.com/delta-ng/Concurrent-Webserver>
- <https://github.com/remzi-arpacidusseau/ostep-projects/tree/master/concurrency-webserver>
- <https://www.petefreitag.com/item/689.cfm>
- https://www.tutorialspoint.com/apache_bench/index.htm